

SI 618

Exploratory Data Analysis

Clustering Analysis

Instructor: Dr. Chris Teplovs (cteplovs@umich.edu)

Lead Developer, Digital Innovation Greenhouse, Office of
Academic Innovation

Adjunct Lecturer of Information, School of Information

GSI: SungJin Nam (sjnam@umich.edu)

Course announcements

- This week:
 - Homework 3 due today
 - Homework 4 (Cluster Analysis) released
- No class next week (Thanksgiving)

» AND.....

This course has been shortened!


- Last day of class is DECEMBER 9, 2016
 - Per information from the SI Registrar
- Due date for project is still DECEMBER 16, 2016
- SLIDES ASSIGNMENT HAS BEEN ELIMINATED
- December 9 class will be brief intro to machine learning, review of 618, and teaching evaluations

SI 618 Data Exploration: Class Schedule

Date	Topic	Assignments Due
Week 1	Course introduction Basics of Programming with R	
Week 2	Basic analysis and visualization using ggplot2: qplot() Manipulating data frames using plyr	Homework 1
Week 3	Smoothing and Trend-finding. Building ggplot Layer by Layer	Homework 2
Week 4	Cluster analysis	Homework 3
Week 5	(Thanksgiving: no class!)	
Week 6	Factor Analysis Methods (PCA, EFA)	Homework 4
Week 7	Machine Learning, Review, Evaluations	

Cluster analysis finds ‘interesting’ groups of objects based on similarity

- What typically makes a ‘good’ clustering?
 - Members are highly similar to each other
 - Minimize within-cluster distances
 - Well-separated from other clusters
 - Maximize between-cluster distances



Applications of Cluster Analysis


- **Understanding**

- Group related documents for browsing
- Group genes and proteins that have similar functionality
- Group stocks with similar price fluctuations

	<i>Discovered Clusters</i>	<i>Industry Group</i>
1	Applied-Matl-DOWN,Bay-Network-Down,3-COM-DOWN,Cabletron-Sys-DOWN,CISCO-DOWN,HP-DOWN,DSC-Comm-DOWN,INTEL-DOWN,LSI-Logic-DOWN,Micron-Tech-DOWN,Texas-Inst-Down,Tellabs-Inc-Down,Natl-Semiconduct-DOWN,Oracl-DOWN,SGI-DOWN,Sun-DOWN	Technology1-DOWN
2	Apple-Comp-DOWN,Autodesk-DOWN,DEC-DOWN,ADV-Micro-Device-DOWN,Andrew-Corp-DOWN,Computer-Assoc-DOWN,Circuit-City-DOWN,Compaq-DOWN,EMC-Corp-DOWN,Gen-Inst-DOWN,Motorola-DOWN,Microsoft-DOWN,Scientific-Atl-DOWN	Technology2-DOWN
3	Fannie-Mae-DOWN,Fed-Home-Loan-DOWN,MBNA-Corp-DOWN,Morgan-Stanley-DOWN	Financial-DOWN
4	Baker-Hughes-UP,Dresser-Inds-UP,Halliburton-HLD-UP,Louisiana-Land-UP,Phillips-Petro-UP,Unocal-UP,Schlumberger-UP	Oil-UP

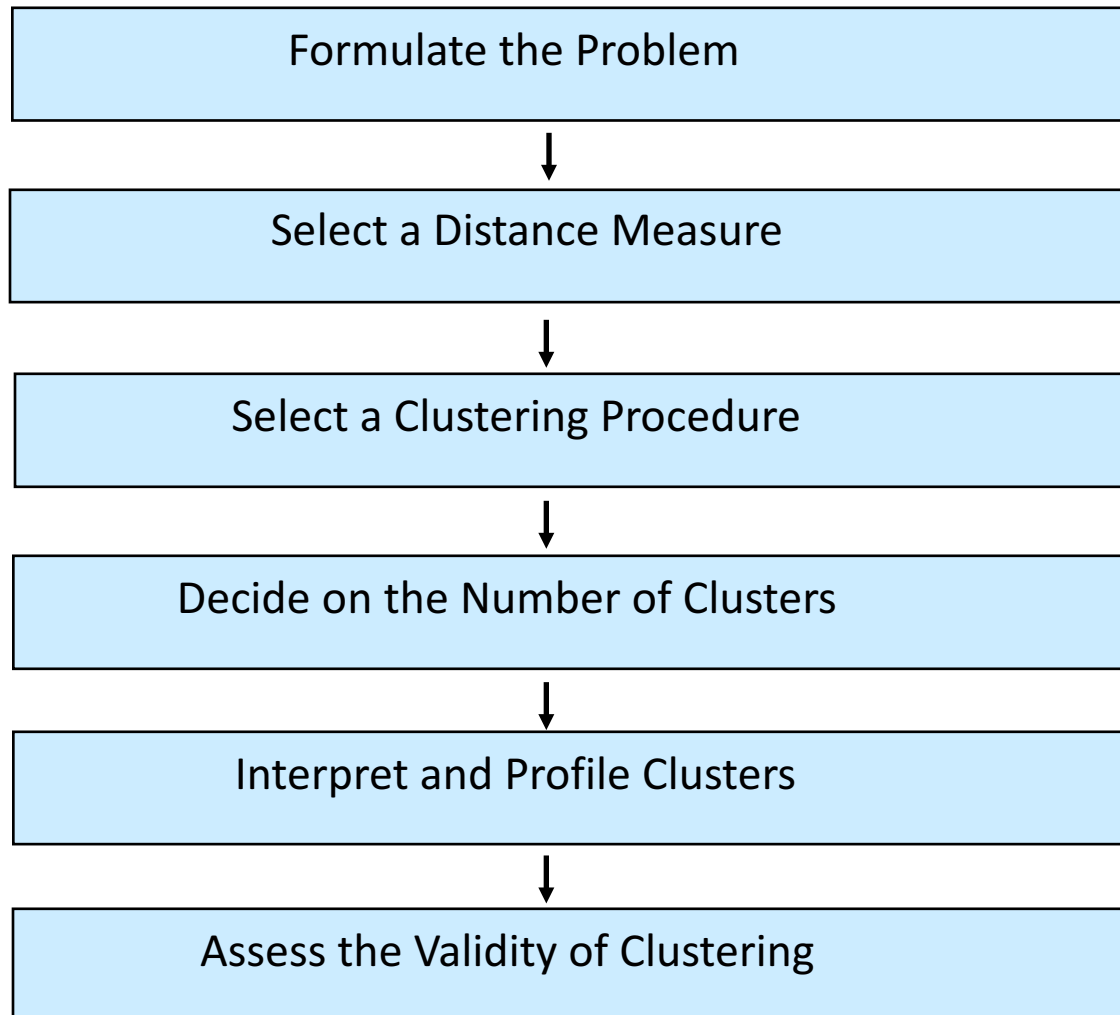
- **Summarization**

- Reduce size of large data sets



Clustering
precipitation in
Australia

Summary: Conducting Cluster Analysis



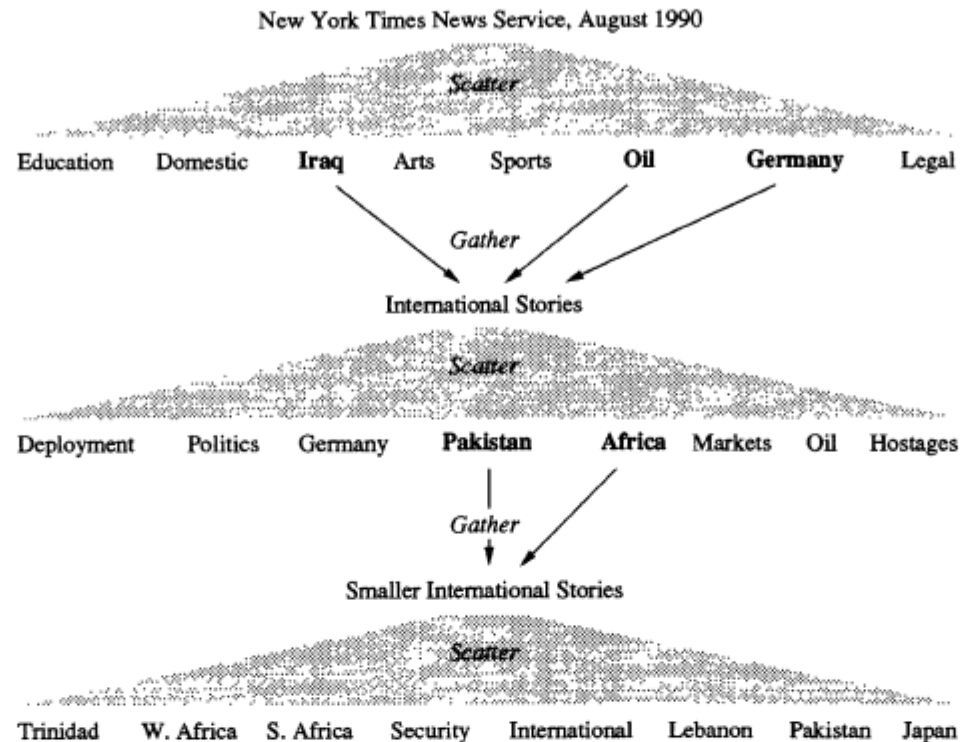
Clustering is often used as an exploratory data analysis tool

- Data understanding
 - Finding underlying factors, groups, structure
- Data navigation
 - Web search and browsing
- Data reduction
 - Clustering creates a new nominal level variable that can be used in any further analysis.
 - In one-dimension, a good way to quantize real-valued variables into k non-uniform buckets
- Data smoothing
 - Infer missing attributes from cluster neighbors

Example: Scatter/Gather. A clustering-based approach to browsing large document collections [Cutting et al. SIGIR 1992]

- What if you have a vague information need that spans topics
- And you're not sure which search terms to use?



1. Scatter: Present the user with a set of clusters
2. Gather: User selects subset of clusters that seem relevant
3. Combine clusters
Repeat from step 1 until done.



Source <http://dl.acm.org/citation.cfm?id=133214>

Example: Clustering lymphoma cancer tissue samples

- B-cell lymphoma go through different stages
- Can we detect which stage automatically?

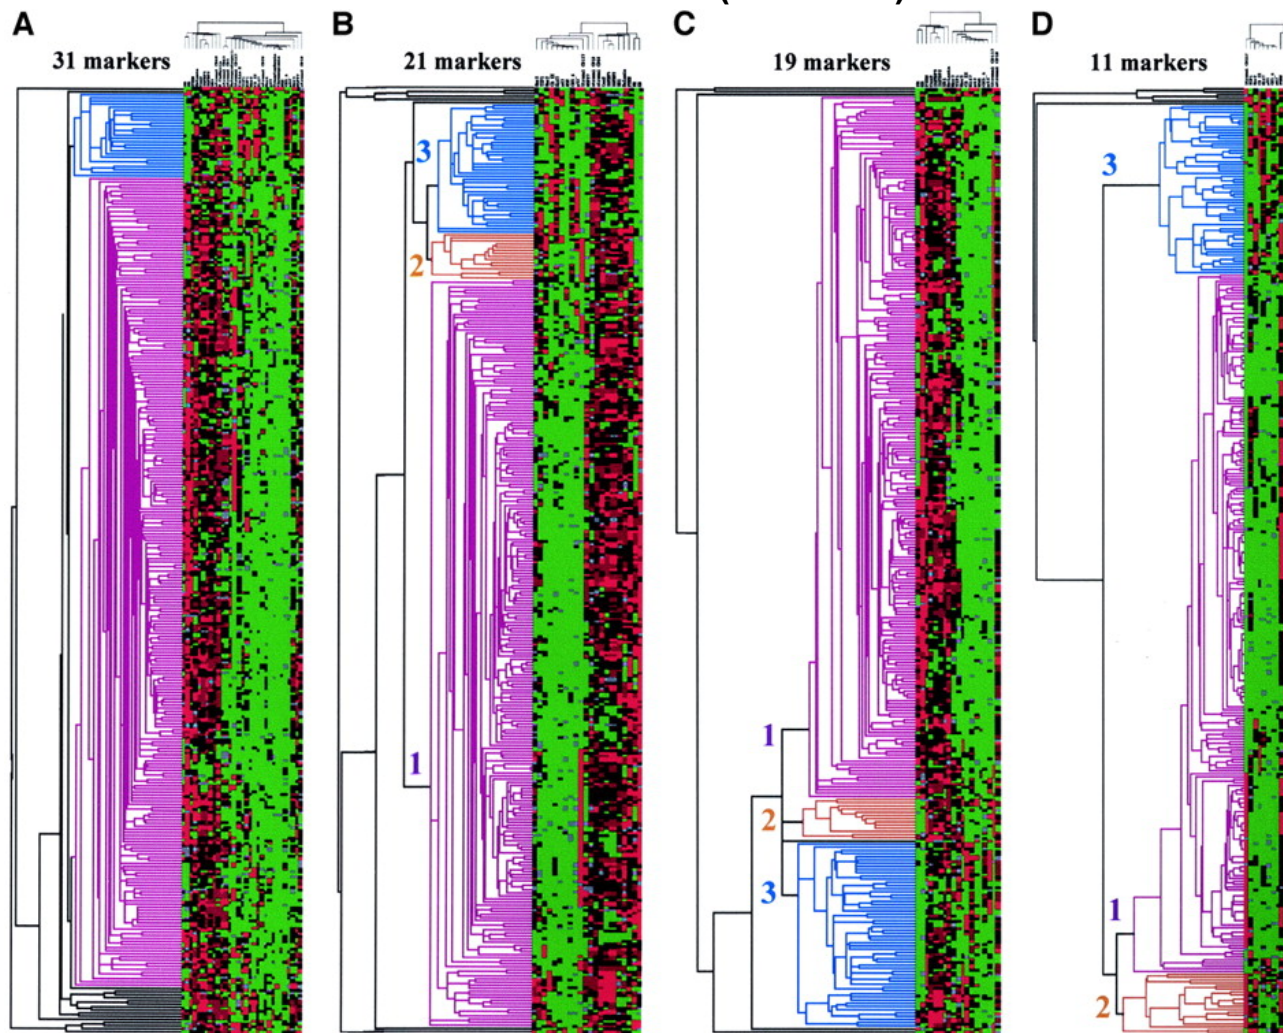


[Source: Chris Ding, ICML 2004 Tutorial on Spectral Clustering]

Clustering arises naturally in many fields


- Health
 - DNA gene expression
 - Cluster cancer variants into treatment groups, based on immunomarkers of cell samples
 - Medical imaging
 - Find likely tumors
- Business
 - Market segments
 - Web site visitors
- Social network analysis
 - Find communities
- Information Retrieval:
 - Search results clustered by similarity, event or topic
 - Personalization for groups of similar users
- Speech understanding
 - Convert waveforms into one of k categories (known as Vector Quantization)

**Hierarchical clustering analysis with 31(A), 21(B), 19(C), and 11(D) immunomarkers.
Groups breast cancer cases (rows) into clinically relevant classes with similar
immunomarkers (columns)**





Makretsov N A et al. Clin Cancer Res 2004;10:6143-6151



Clustering can be ambiguous: What is the 'best' clustering here?





How many clusters?




Six Clusters



Two Clusters



Four Clusters





What algorithms are used to find clusters?

Answer: huge range of approaches

- Assigning objects to clusters
 - ‘Hard’ (partitional) each object belongs to exactly 1 cluster
 - ‘Soft’ : each object can belong to multiple clusters
- Hierarchical vs non-hierarchical
 - A set of nested clusters organized as a tree
- By far most widely-used fall into two types:
 - **Heirarchical**: agglomerative, single-link, etc.
 - **Partitional**: k-means, k-median, etc.


Hierarchical Clustering

- Produces a set of nested clusters organized as a hierarchical tree
- Can be visualized as a dendrogram
 - A tree like diagram that records the sequences of merges or splits



Strengths of Hierarchical Clustering

- Do not have to assume any particular number of clusters
 - Any desired number of clusters can be obtained by ‘cutting’ the dendrogram at the proper level
- They may correspond to meaningful taxonomies
 - Example in biological sciences (e.g., animal kingdom, phylogeny reconstruction, ...)




Hierarchical clustering

- Bottom-up ('Agglomerative')
 - Start with each point being in its own cluster
 - At each step
 - Merge the most similar pair of clusters based on a cost function
 - Continue until you have k clusters, or everything is in one big cluster
- Top-down ('Divisive')
 - Start with all points in a single big cluster
 - At each step:
 - Split the cluster into two smaller clusters based on a cost function
 - Continue until you have k clusters, or each point is in its own cluster

http://wiki.stat.ucla.edu/socr/index.php/SOCR_EduMaterials_AnalysisActivities_HierarchicalClustering

Agglomerative (bottom-up) Clustering: Starting Situation

- Start with clusters of individual points and a proximity matrix of object-to-object distances



	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix




Agglomerative Clustering Algorithm

- More popular hierarchical clustering technique
- Basic algorithm is straightforward
 1. Compute the proximity matrix
 2. Let each data point be a cluster
 3. **Repeat**
 4. Merge the two closest clusters
 5. Update the proximity matrix
 6. **Until** only a single cluster remains
- Key operation: computation of the proximity of two clusters. The cost function.
 - Different approaches to defining the distance between clusters distinguish the different algorithms

Agglomerative (bottom-up) Clustering: Starting Situation

- Start with clusters of individual points and a proximity matrix of object-to-object distances




	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix




Intermediate Situation

- After some merging steps, we have some clusters




	C1	C2	C3	C4	C5
C1					
C2					
C3					
C4					
C5					

Proximity Matrix



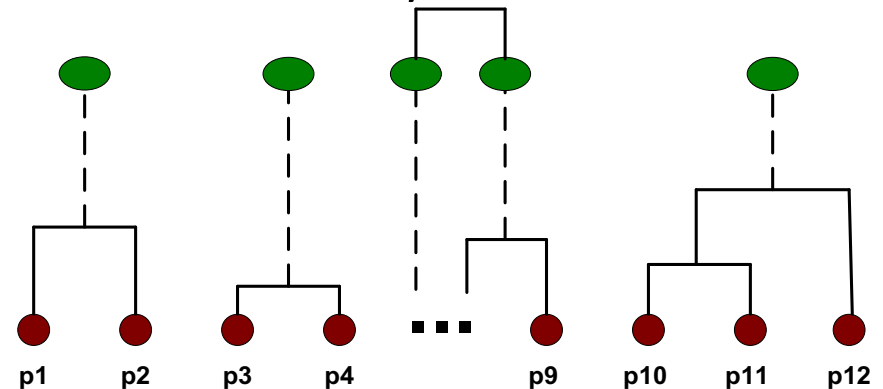
Intermediate Situation

- We want to merge the two closest clusters (C2 and C5) and update the proximity matrix.




	C1	C2	C3	C4	C5
C1					
C2					
C3					
C4					
C5					

Proximity Matrix




After Merging

- The question is “How do we update the proximity matrix?”




	C1	C2 + C5	C3	C4
C1		?		
C2 + C5	?	?	?	?
C3		?		
C4		?		

Proximity Matrix



How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
 - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
 - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
 - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
 - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity






- MIN
- MAX
- Group Average
- Distance Between Centroids
- Other methods driven by an objective function
 - Ward's Method uses squared error

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						


Proximity Matrix

Cost functions for bottom-up (agglomerative) clustering

- Single linkage
 - Minimum distance between clusters
- Complete linkage
 - Max distance between clusters
- Average linkage
 - Average distance between clusters




Cluster Similarity: MIN or Single Linkage




- Similarity of two clusters is based on the two most similar (closest) points in the different clusters
 - Determined by one pair of points, i.e., by one link in the proximity graph.


	I1	I2	I3	I4	I5
I1	1.00	0.90	0.10	0.65	0.20
I2	0.90	1.00	0.70	0.60	0.50
I3	0.10	0.70	1.00	0.40	0.30
I4	0.65	0.60	0.40	1.00	0.80
I5	0.20	0.50	0.30	0.80	1.00



Hierarchical Clustering: MIN




Nested Clusters




Dendrogram

Strength of MIN




Original Points




Two Clusters

- Can handle non-elliptical shapes

Limitations of MIN




Original Points



Two Clusters


- Sensitive to noise and outliers

Cluster Similarity: MAX or Complete Linkage




- Similarity of two clusters is based on the two least similar (most distant) points in the different clusters
 - Determined by all pairs of points in the two clusters


	I1	I2	I3	I4	I5
I1	1.00	0.90	0.10	0.65	0.20
I2	0.90	1.00	0.70	0.60	0.50
I3	0.10	0.70	1.00	0.40	0.30
I4	0.65	0.60	0.40	1.00	0.80
I5	0.20	0.50	0.30	0.80	1.00



Hierarchical Clustering: MAX




Nested Clusters




Dendrogram

Strength of MAX




Original Points




Two Clusters

- Less susceptible to noise and outliers

Limitations of MAX



Original Points




Two Clusters

- Tends to break large clusters
- Biased towards globular clusters

Cluster Similarity: Group Average


- Proximity of two clusters is the average of pairwise proximity between points in the two clusters.




$$\text{proximity}(\text{Cluster}_i, \text{Cluster}_j) = \frac{\sum_{\substack{p_i \in \text{Cluster}_i \\ p_j \in \text{Cluster}_j}} \text{proximity}(p_i, p_j)}{|\text{Cluster}_i| * |\text{Cluster}_j|}$$

- Need to use average connectivity for scalability since total proximity favors large clusters


	I1	I2	I3	I4	I5
I1	1.00	0.90	0.10	0.65	0.20
I2	0.90	1.00	0.70	0.60	0.50
I3	0.10	0.70	1.00	0.40	0.30
I4	0.65	0.60	0.40	1.00	0.80
I5	0.20	0.50	0.30	0.80	1.00



Hierarchical Clustering: Group Average



Nested Clusters



Dendrogram

Hierarchical Clustering: Group Average

- Compromise between Single and Complete Link
- Strengths
 - Less susceptible to noise and outliers
- Limitations
 - Biased towards globular clusters

Ward's method (1963)

- **Ward's distance** between clusters C_i and C_j is the *difference* between the *total within cluster sum of squares for the two clusters separately*, and the *within cluster sum of squares resulting from merging the two clusters* in cluster C_{ij}


$$D_w(C_i, C_j) = \sum_{x \in C_i} (x - r_i)^2 + \sum_{x \in C_j} (x - r_j)^2 - \sum_{x \in C_{ij}} (x - r_{ij})^2$$

- r_i : centroid of C_i
- r_j : centroid of C_j
- r_{ij} : centroid of C_{ij}


Ward's distance for clusters

- Similar to group average and centroid distance
- Less susceptible to noise and outliers
- Biased towards globular clusters
- Hierarchical analogue of k-means
 - Can be used to initialize k-means


Hierarchical Clustering: Comparison



MIN




MAX



Group Average

Ward's Method



Which type of hierarchical clustering to use?


- Different methods have different strengths and weaknesses.
 - Ward's method tends to give equal sized clusters
 - Single linkage (nearest neighbor) tends to make long strings into a cluster.
 - Top-down is sensitive to early errors: bad first choice can wreck the entire process
 - Bottom-up can't see the whole dataset

- Two major clustering algorithms
 - Hierarchical
 - K-means
- General questions:
 - How many clusters is best?
 - How can we assess and visualize cluster quality?
 - How can we visualize clusters?

K-means: the other massively popular clustering method.. and very different in nature

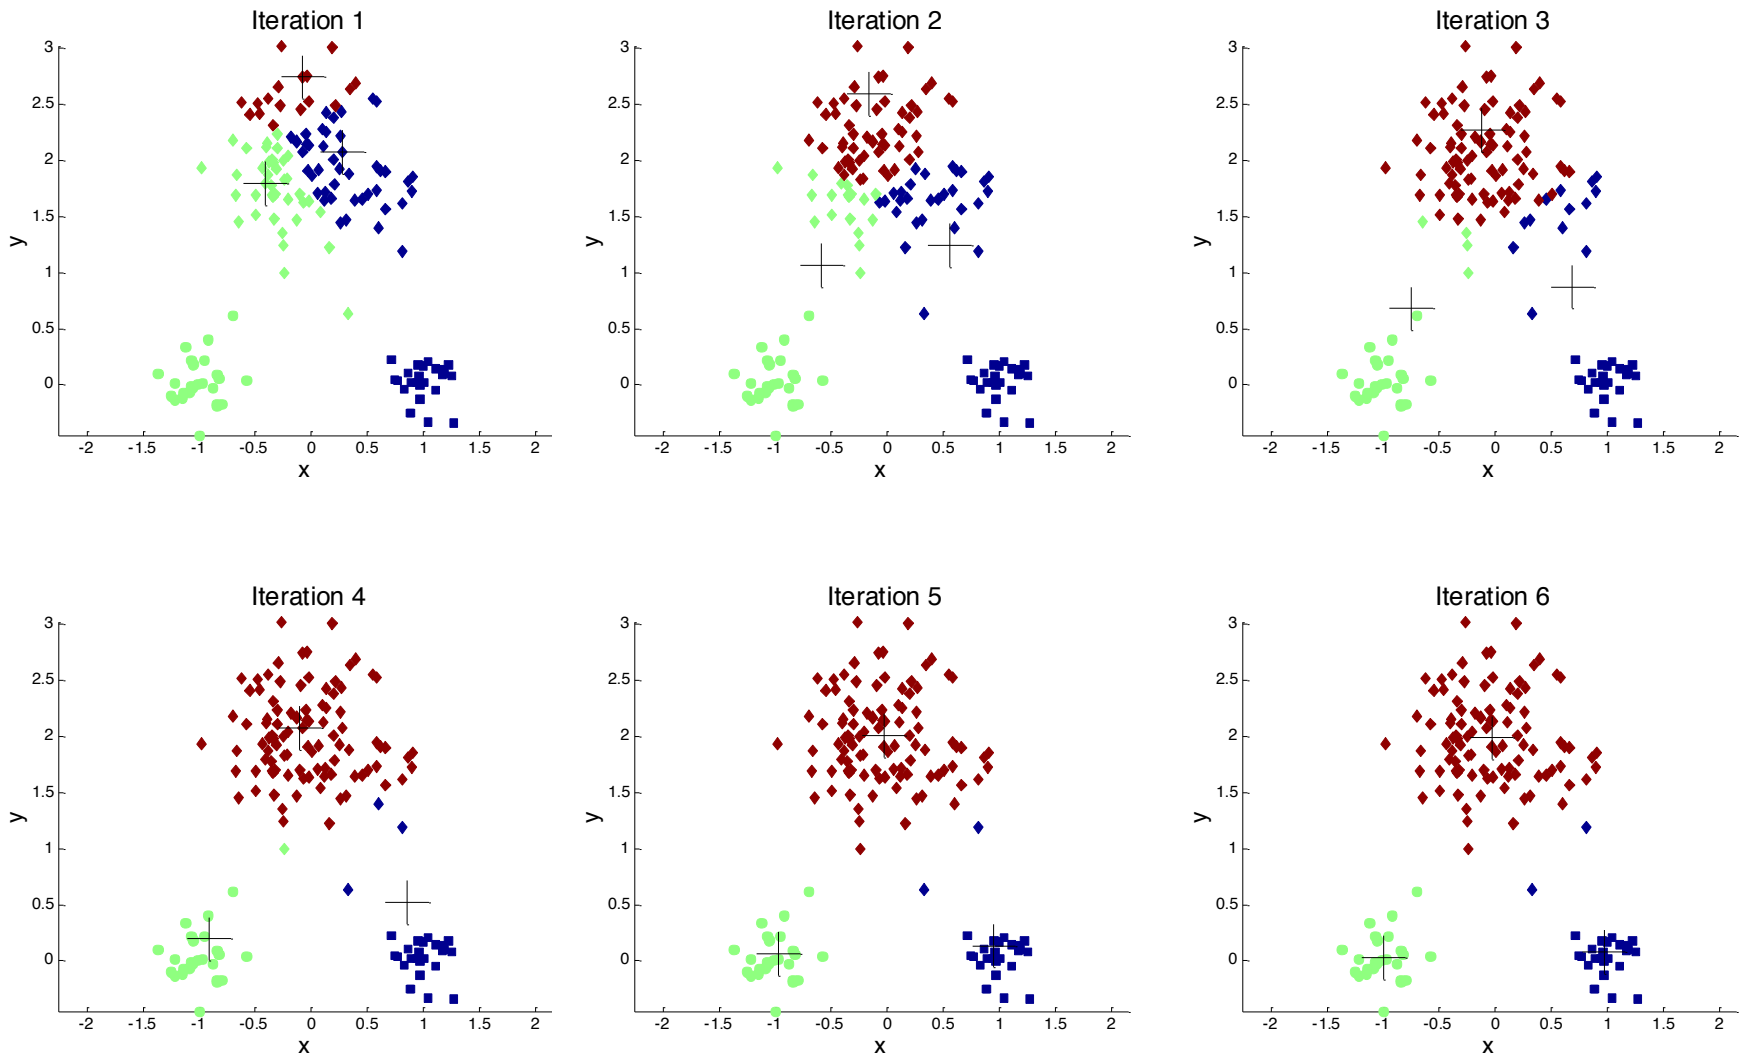
- Partitional clustering approach
- Each cluster associated with a **centroid** (centerpoint)
- Each point is assigned to the cluster with the closest centroid
- Number of clusters, K , must be specified in advance
- The basic algorithm is very simple

-
- 1: Select K points as the initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning all points to the closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** The centroids don't change
-



<http://stat.ethz.ch/R-manual/R-devel/library/stats/html/kmeans.html>

The k-means algorithm ($k = 3$)





K-means is a special case of model-based clustering

- Assume data generated from **k** probability distributions
- **Goal:** find the distribution parameters
- **Algorithm:** Expectation Maximization (EM)
- **Output:** Distribution parameters and a **soft** assignment of points to clusters


K-means Clustering – Details

- Different initializations can result in different solutions
 - Initial centroids are often chosen randomly.
 - Clusters produced vary from one run to another.
 - So multiple runs are sometimes done
- Centroid is typically the mean of the points in the cluster.
 - K-medoid: center must be an actual datapoint. Useful when mean of a feature is not defined or available
- ‘Closeness’ is measured by Euclidean distance, cosine similarity, correlation, etc.
- K-means will converge for common similarity measures mentioned above.
- Most convergence happens in the first few iterations.
 - Often the stopping condition is changed to ‘Until relatively few points change clusters’
- Complexity is $O(n * K * l * d)$
 - n = number of points, K = number of clusters,
 l = number of iterations, d = number of attributes

Two different K-means Clusterings




Optimal Clustering? No.




Optimal Clustering






Importance of Choosing Initial Centroids



Importance of Choosing Initial Centroids ...



Importance of Choosing Initial Centroids ...



Trying to find good optimal k-means clusterings

- Idea 1: Be careful about where you start
 - Place first center on randomly chosen datapoint
 - Place second centroid on datapoint as far as possible from first center (or soft probabilistic version thereof)
 - Place j -th center on datapoint that's as far as possible from centers 1 thru $j - 1$
- Idea 2: Do many runs of k-means
 - Each from a different random start configuration
- Many heuristics around

Limitations of K-means

- K-means has problems when clusters are of differing
 - Sizes
 - Densities
 - Non-globular shapes
- K-means has problems when the data contains outliers.

When should you use k-means vs hierarchical approach?

- Do you need to easily interpret the clusters?
- Do you know the right k ?
- Hierarchical clustering is the sort that you might apply when there is a "tree" structure to the data (e.g. living things).
- K-means clustering does not assume a tree structure.
- If you have only two or three dimensions (or can sensibly reduce your data by factor analysis) you can plot it and see what sort of relationships you have. Are you looking for nice spherical clusters, or are long chains more suitable?
- k-means prefers solutions where clusters are of similar size
 - very different cluster sizes, shapes, densities can confuse it
 - complex cluster geometry, or outliers
 - need to specify and test for good k choice
- Can combine the two approaches, e.g.
 1. Try several hierarchical methods and see which gives the most interpretable clusters.
 2. Use k-means (with the hierarchical cluster centroids as starting points) to clean up the hierarchical cluster.

Clustering in R

Step 1: Data preparation

```
# Prepare Data  
mydata <- na.omit(mydata) # listwise deletion of missing  
mydata <- scale(mydata) # standardize variable scales
```

Note: Scaling is important. Think about what happens if points are clustered on one variable from 0-100 and another on 0.0-1.0

Reference: <http://www.statmethods.net/advstats/cluster.html>

Step 2: Clustering (if Hierarchical)

```
> head(cars.data)
```

	MPG	Weight	Drive_Ratio	Horsepower	Displacement	Cylinders
Buick Estate Wagon	16.9	4.360	2.73	155	350	8
Ford Country Squire Wagon	15.5	4.054	2.26	142	351	8
Chevy Malibu Wagon	19.2	3.605	2.56	125	267	8
Chrysler LeBaron Wagon	18.5	3.940	2.45	150	360	8
Chevette	30.0	2.155	3.70	68	98	4
Toyota Corona	27.5	2.560	3.05	95	134	4

```
# Heirarchical clustering: compute distance matrix
```

```
cars.dist = dist(cars.data)
```

```
> as.matrix(cars.dist)
```

	Buick Estate Wagon	Ford Country Squire Wagon	Chevy Malibu Wagon	Chrysler LeBaron Wagon	Chevette	Toyota Corona
Buick Estate Wagon	0.00000	13.125339	88.28867	11.30552	266.9576988	224.472053
Ford Country Squire Wagon	13.12534	0.000000	85.78451	12.41165	264.0396368	222.397968
Chevy Malibu Wagon	88.28867	85.784507	0.00000	96.30480	178.7345577	136.657316
Chrysler LeBaron Wagon	11.30552	12.411652	96.30480	0.00000	274.8108417	232.809502
Chevette	266.95770	264.039637	178.73456	274.81084	0.0000000	45.075897
Toyota Corona	224.47205	222.397968	136.65732	232.80950	45.0758974	0.000000

```
cars.hclust <- hclust(cars.dist, method = "average")
```

Reference: <http://www.statmethods.net/advstats/cluster.html>

Step 2: Partitioning (if k-means)

```
# K-Means Cluster Analysis
> fit <- kmeans(cars.data, 5) # 5 cluster solution

> fit
```

K-means clustering with 5 clusters of sizes 10, 4, 6, 11, 7

Cluster means:

	MPG	Weight	Drive_Ratio	Horsepower	Displacement	Cylinders
1	25.59000	2.638100	3.298000	96.50000	133.50000	4.3
2	19.12500	3.503750	2.682500	115.00000	245.25000	6.5
3	21.91667	2.970833	3.128333	113.66667	173.83333	6.0
4	32.43636	2.078636	3.477273	70.90909	94.63636	4.0
5	17.17143	3.957714	2.402857	139.85714	333.85714	8.0

Clustering vector:

Buick Estate Wagon	Ford Country Squire Wagon	Chevy Malibu Wagon	Chrysler LeBaron Wagon	Chevette
5	5	2	5	4
Toyota Corona	Datsun 510	Dodge Omni	Audi 5000	Volvo 240 GL
1	1	4	1	3
Saab 99 GLE	Peugeot 694 SL	Buick Century Special	Mercury Zephyr	Dodge Aspen
1	3	2	3	2
AMC Concord D/L	Chevy Caprice Classic	Ford LTD	Mercury Grand Marquis	Dodge St Regis
2	5	5	5	5
Ford Mustang 4	Ford Mustang Ghia	Mazda GLC	Dodge Colt	AMC Spirit
1	3	4	4	1
VW Scirocco	Honda Accord LX	Buick Skylark	Chevy Citation	Olds Omega
4	4	1	3	3
Pontiac Phoenix	Plymouth Horizon	Datsun 210	Fiat Strada	VW Dasher
1	4	4	4	4
Datsun 810	BMW 320i	VW Rabbit		
1	1	4		


```
# get cluster means
aggregate(mydata, by=list(fit$cluster), FUN=mean)
```

```
# append cluster assignment
mydata <- data.frame(mydata, fit$cluster)
```

Reference: <http://www.statmethods.net/advstats/cluster.html>

Step 3: Visualizing

```
plot(cars.hclust, labels=cars$Car, main='Hierarchical cluster analysis using average linkage')
```




Reference: <http://www.statmethods.net/advstats/cluster.html>

Old Clustering in R: heatmap

```
> mtscaled <- as.matrix(scale(mtcars))  
> heatmap(mtscaled, Colv=F,  
scale='none')
```



Clustering columns:

- Some info is highly correlated
- e.g. displacement, hp, # cylinders very similar



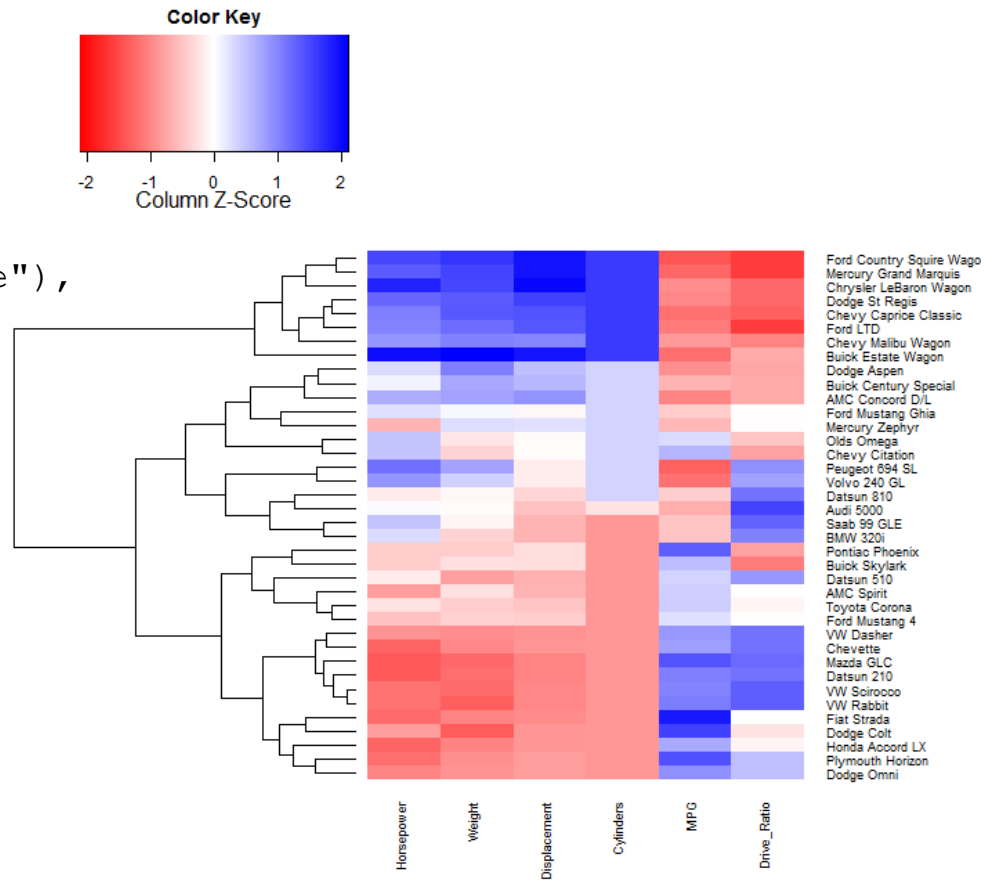
New clustering in R: heatmap.2

```
> mtscaled <-  
as.matrix(scale(mtcars))  
> heatmap.2(mtscaled,  
Colv=F, scale='none')
```



Clustering in R: heatmap.2

```
install.packages("gplots")
library(gplots)
heatmap.2(as.matrix(cars.data),
hclustfun = function(x)
      hclust(x,method = "average"),
scale = "column",
dendrogram="row",
trace="none",
density.info="none",
col=redblue(256),
lhei=c(2,5.0), lwid=c(1.5,2.5),
keysize = 0.25,
margins = c(5, 8),
cexRow=0.7,cexCol=0.7)
```



Clustering in R: heatmap.2

Reference: <http://cran.r-project.org/web/packages/gplots/gplots.pdf>

heatmap.2 (x,

```
# dendrogram control
Rowv = TRUE,
Colv=if(symm)"Rowv" else TRUE,
distfun = dist,
hclustfun = hclust,
dendrogram = c("both", "row", "column", "none"),
symm = FALSE,

# data scaling
scale = c("none", "row", "column"),
na.rm=TRUE,

# image plot
revC = identical(Colv, "Rowv"),
add.expr,

# mapping data to colors
breaks,
symbreaks=min(x < 0, na.rm=TRUE) || scale!="none",

# colors
col="heat.colors",

# block separation
colsep,
rowsep,
sepcolor="white",
sepwidth=c(0.05,0.05),

# cell labeling
cellnote,
notecex=1.0,
notecol="cyan",
na.color=par("bg"),

# level trace
trace=c("column", "row", "both", "none"),
tracecol="cyan",
hline=median(breaks),
vline=median(breaks),
linecol=tracecol,

# Row/Column Labeling
margins = c(5, 5),
ColSideColors,
RowSideColors,
cexRow = 0.2 + 1/log10(nr),
cexCol = 0.2 + 1/log10(nc),
labRow = NULL,
labCol = NULL,

# color key + density info
key = TRUE,
keysize = 1.5,
density.info=c("histogram", "density", "none"),
denscol=tracecol,
symkey = min(x < 0, na.rm=TRUE) || symbreaks,
densadj = 0.25,

# plot labels
main = NULL,
xlab = NULL,
ylab = NULL,

# plot layout
lmat = NULL,
lhei = NULL,
lwid = NULL,
```



How many clusters?

- Theoretical, conceptual or practical issues may suggest a certain number of clusters
- Hierarchical clustering:
 - Distance threshold at which clusters are combined
- K-means and other non-hierarchical
 - Ratio of total within-groups variance to between-group variance, vs # of clusters
 - Special case: within-groups sum of squares vs # of clusters
 - Elbow/sharp bend shows point at which adding more clusters helps reduce distortion measure less and less

How many clusters?

```
# Determine number of clusters
wss <- (nrow(mydata)-1)*sum(apply(mydata,2,var))
for (i in 2:15) wss[i] <- sum(kmeans(mydata,
  centers=i)$withinss)
plot(1:15, wss, type="b", xlab="Number of Clusters",
  ylab="Within groups sum of squares")
```

How many clusters?




How do we know if we've found good quality clusters?

- Compare cluster stability across:
 - Different distance measures
 - Different clustering methods
 - Different 50/50 random data splits
 - Different variable/features deletions
 - Different data orderings (non-hierarchical)
- “Good” clusterings (if they exist) are generally stable and robust to perturbations in methods or data

Silhouette scores

[Peter J. Rousseeuw](#) (1987). "Silhouettes: a Graphical Aid to the Interpretation and Validation of Cluster Analysis". *Computational and Applied Mathematics* **20**: 53–65.

- A graphical aid for interpretation and validation of cluster analysis
- $a(i)$: average dissim of datum i with others in same cluster
- $b(i)$: lowest average dissim for other clusters (neighboring cluster)
- Gives degree of confidence in cluster assignment
 - Well-clustered elements: score near 1
 - Poorly-clustered elements: score near -1 (probably in wrong cluster)
- To compute in R: `silhouette(cars.pam, cars.dist)`



$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$


Which can be written as:

$$s(i) = \begin{cases} 1 - a(i)/b(i), & \text{if } a(i) < b(i) \\ 0, & \text{if } a(i) = b(i) \\ b(i)/a(i) - 1, & \text{if } a(i) > b(i) \end{cases}$$

See `cluster_analysis.R` example file

Spectral clustering


- How does k-means or hierarchical clustering deal with THIS?



- Not well: data has obvious local structure (lies along curves) but traditional clustering methods don't account for that.

Spectral clustering


1. Map points in regular feature space to 'spectral' space
2. Then apply conventional clustering
 - e.g. k-means




Source: Chris Ding, ICML 2004 Tutorial on Spectral Clustering

Spectral clustering methods work well for data with local geometry structure, i.e. points tend to form curves or surfaces

In case you were curious: The key idea of spectral clustering is to approximate the optimal cluster indicator function by the second eigenvector of the well-known graph Laplacian. It amounts to finding the optimal balanced cut of an undirected weighted graph where the weights represent the similarities between points.



k-means

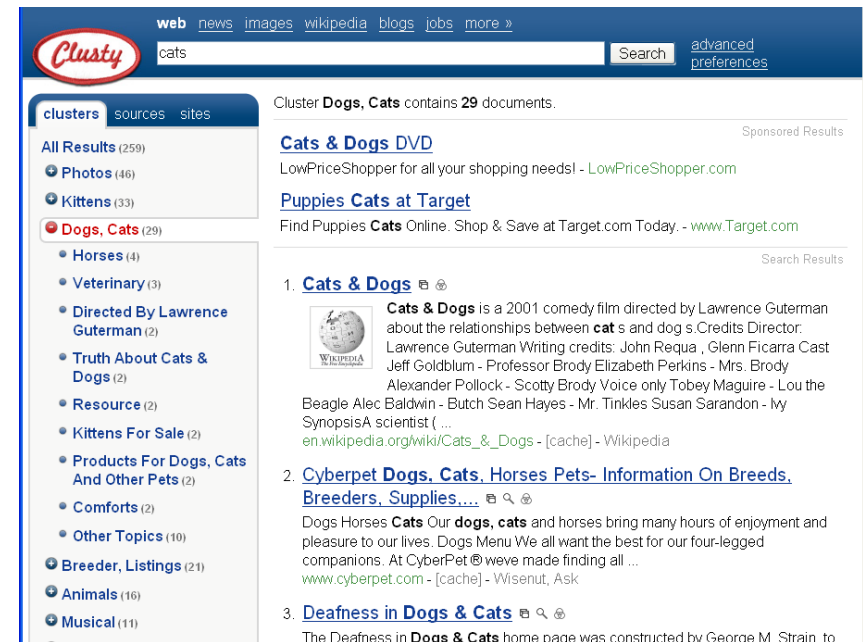


Spectral clustering

Source: <http://www.ml.uni-saarland.de/code/pSpectralClustering/pSpectralClustering.html>

How to automatically name or label clusters?


- What determines 'good' vs 'bad' names?
 - Usually task-based evaluations, e.g. search interface effect on user satisfaction/effectiveness
- Use a text summary of a representative element
 - Centroid, medoid, etc.
- Classify with existing hierarchy
- Don't show labels
 - Benefits unclear depending on scenario



See: http://searchuserinterfaces.com/book/sui_ch8_navigation_and_search.html

How to automatically name or label clusters?

- Use crowd-powered methods to create and label object hierarchies [Chilton et al, CHI 2013]
 - And more generally, global pictures of a dataset



Source: <http://hmslydia.com/cascade.html>

How can we visualize high-dimensional clusters?
One method: multi-dimensional scaling (MDS)

- “Flatten” multidimensional cloud to 2d and preserve distances as much as possible
- Input
 - List of N datapoints (m -dimensional)
 - Or: derived $N \times N$ distance matrix
- Output: list of 2-dimensional datapoints
- Preserve the true relative distances between all pairs of m -dimensional points

MDS and PAM in R

Partitioning Around Medoids (find k representative objects, then cluster)

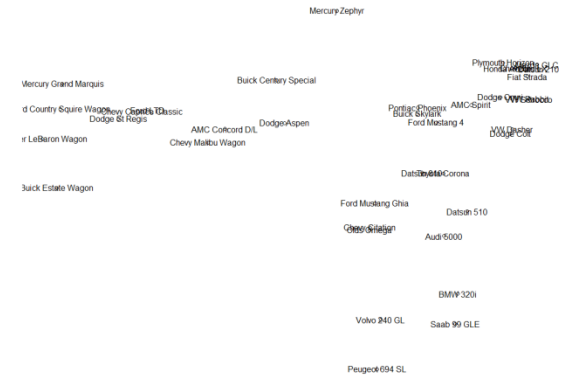
- Extract feature points (rows) for objects
- Compute d : matrix of distances between rows
- Call `cmdscale(d)`
- Plot 2-d
- ```
Classical MDS
N rows (objects) x p columns (variables)
each row identified by a unique row name

d <- dist(cars.data) # euclidean distances
 between the rows
fit <- cmdscale(d,eig=TRUE, k=2) # k is the
 number of dim
fit # view results


plot solution
x <- fit$points[,1]
y <- fit$points[,2]
plot(x, y, xlab="Coordinate 1",
 ylab="Coordinate 2",
 main="Metric MDS", type="n")
text(x, y, labels = rownames(fit$points))
```
- Also could use
 

```
cars.pam = pam(cars.dist, 3)
clusplot(cars.pam, labels=2)
```

### MDS

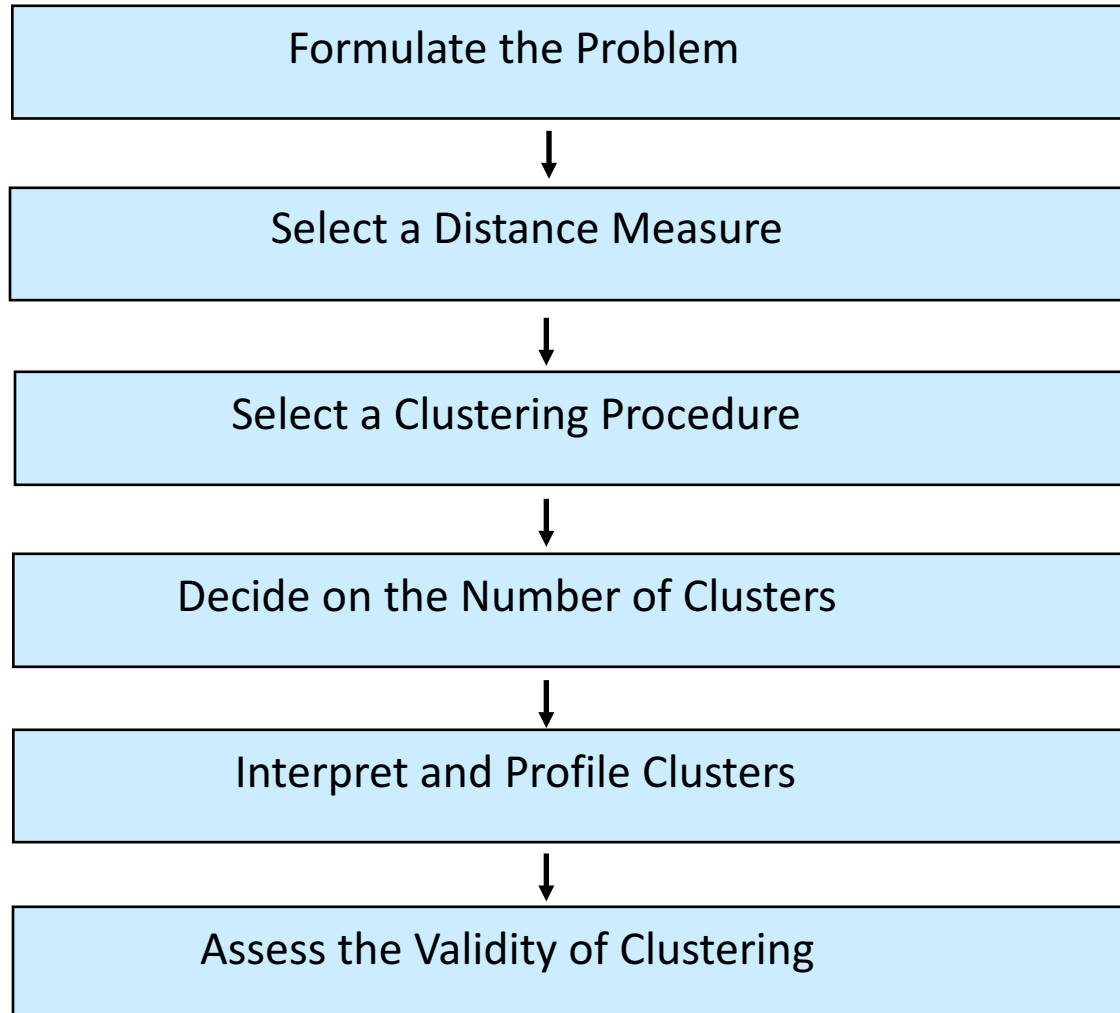


### PAM



<http://stat.ethz.ch/R-manual/R-devel/library/cluster/html/pam.html>

# Summary: Conducting Cluster Analysis



# Homework 4: Cluster Analysis



## Cars

- Cluster analysis of a set of 38 vehicles
- Each row in the data set contains various types of information about each vehicle

### What to do:

- Use hierarchical and k-means clustering to cluster car data
- Visualize the results in various ways discussed in the lecture

|    | Country | Car                       | MPG  | Weight | Drive_Ratio | Horsepower | Displacement | cylinders |
|----|---------|---------------------------|------|--------|-------------|------------|--------------|-----------|
| 1  | U.S.    | Buick Estate wagon        | 16.9 | 4.360  | 2.73        | 155        | 350          | 8         |
| 2  | U.S.    | Ford Country Squire wagon | 15.5 | 4.054  | 2.26        | 142        | 351          | 8         |
| 3  | U.S.    | chevy Malibu wagon        | 19.2 | 3.605  | 2.56        | 125        | 267          | 8         |
| 4  | U.S.    | chrysler LeBaron wagon    | 18.5 | 3.940  | 2.45        | 150        | 360          | 8         |
| 5  | U.S.    | chevette                  | 30.0 | 2.155  | 3.70        | 68         | 98           | 4         |
| 6  | Japan   | Toyota Corona             | 27.5 | 2.560  | 3.05        | 95         | 134          | 4         |
| 7  | Japan   | Datsun 510                | 27.2 | 2.300  | 3.54        | 97         | 119          | 4         |
| 8  | U.S.    | Dodge omni                | 30.9 | 2.230  | 3.37        | 75         | 105          | 4         |
| 9  | Germany | Audi 5000                 | 20.3 | 2.830  | 3.90        | 103        | 131          | 5         |
| 10 | Sweden  | volvo 240 GL              | 17.0 | 3.140  | 3.50        | 125        | 163          | 6         |



# What you should know

- Basic use of hierarchical and k-means clustering
- When is k-means clustering preferable to hierarchical clustering? Or vice-versa?
- How is cluster quality measured?
- How to do clustering in R